

1. Ridge Regression :

A. The ridge regression problem is : $\min_{\beta} ||X\beta - y||^2 + \lambda ||\beta||^2$

By solving it, the closed form solution is

$$\beta = (\lambda I + X^T X)^{-1} X^T y$$

Augmented data of

Matrix $X' = \begin{pmatrix} X \\ \sqrt{\lambda} I \end{pmatrix}$ and response vector $y' = \begin{pmatrix} y \\ 0_k \end{pmatrix}$

Where I is $k \times k$ identity matrix and 0_k is $k \times 1$ vector. Calculating $X'^T X'$

$$X'^T X' = (X^T \sqrt{\lambda} I) \cdot \begin{pmatrix} X \\ \sqrt{\lambda} I \end{pmatrix} = X^T X + \lambda I$$

And $X'^T y$ can be calculated as

$$X'^T y = (X^T \sqrt{\lambda} I) \cdot \begin{pmatrix} y \\ 0_k \end{pmatrix} = X^T y$$

We can calculate the least squares regression closed solution

$$\begin{aligned} \hat{\beta} &= (X'^T X')^{-1} X'^T y' \\ \Rightarrow \hat{\beta} &= (X^T X + \lambda I)^{-1} X^T y \end{aligned}$$

Which is the closed form solution to ridge regression with parameter λ . We proved that the ridge regression estimates can be obtained by ordinary least squares regression on an augmented data set.

2a) Preprocessing the data.

```
def preprocess_data(trainx, valx, testx):  
    train_mean = np.mean(trainx, axis=0)  
    train_std = np.std(trainx, axis=0)  
    trainx = (trainx - train_mean) / train_std  
    valx = (valx - train_mean) / train_std  
    testx = (testx - train_mean) / train_std  
    return trainx, valx, testx
```

Initially, I scaled the features to have similar values by standardizing it (calculated the mean and standard deviation of each feature on the training data, then standardized the training, validation, and test data using those statistics.) This can also be achieved using standard scalar but I preferred this method. I chose to standardize because it prevents large value features from dominating. It makes all the features equally important and can improve the performance of the model.

Later on I did some feature engineering to add some features and removed the outliers to improve the accuracy of the model.

```
In [3]: df["hours"]=(df['date'].str.split(':').str[0].str.split(" ").str[1]).astype(str).astype(int)
df['hour*lights'] = df.hours * df.lights
df['hour_avg'] = list(map(dict(df.groupby('hours')['Appliances'].mean()).get, df.hours))
df_val["hours"]=(df_val['date'].str.split(':').str[0].str.split(" ").str[1]).astype(str).astype(int)
df_val['hour*lights'] = df_val.hours * df_val.lights
df_val['hour_avg'] = list(map(dict(df_val.groupby('hours')['Appliances'].mean()).get, df_val.hours))
df_test["hours"]=(df_test['date'].str.split(':').str[0].str.split(" ").str[1]).astype(str).astype(int)
df_test['hour*lights'] = df_test.hours * df_test.lights
df_test['hour_avg'] = list(map(dict(df_test.groupby('hours')['Appliances'].mean()).get, df_test.hours))

In [6]: feature_set = ['low_consum', 'high_consum', 'hours', 't6', 'rh_6', 'lights', 'hour*lights',
                       'tdewpoint', 'visibility', 'press_mm_hg', 'windspeed']

In [37]: corr_matrix_train = df.corr(method='pearson')
matrix_train=corr_matrix_train["Appliances"].sort_values(ascending=False)
print((matrix_train))
corr_matrix_val = df_val.corr(method='pearson')
matrix_val=corr_matrix_val["Appliances"].sort_values(ascending=False)
print((matrix_val))
corr_matrix_test = df_test.corr(method='pearson')
matrix_test=corr_matrix_test["Appliances"].sort_values(ascending=False)
print((matrix_test))
```

```
: X_train=df.drop(['Appliances', 'date'], axis = 1).values
Y_train = df[['Appliances']].values
X_val=df_val.drop(['Appliances', 'date'], axis = 1).values
Y_val = df_val[['Appliances']].values
X_test=df_test.drop(['Appliances', 'date'], axis = 1).values
Y_test = df_test[['Appliances']].values
train_mean = np.mean(X_train, axis=0)
train_std = np.std(X_train, axis=0)
X_train = (X_train - train_mean) / train_std
X_val = (X_val - train_mean) / train_std
X_test = (X_test - train_mean) / train_std
```

2e) Comparing the RMSE and R^2 of 2c and 2d

Model	Train RMSE	Train R^2	Val RMSE	Val R^2	Test RMSE	Test R^2
Train only (2c)	94.164	0.25269	91.610	0.12024	88.04	0.0611
Train + val (2d)	95.021	0.2390	85.0001	0.24262	82.696	0.1718

Result :

The model trained on training + validation data(2d) performs better on test data rather than only trained on training data(2c). This is because 2d is trained with more data. But the train metrics decrease when trained with both the data. We can also see a slight increase in RMSE in train metrics. Coming to the Validation metrics, we can see a slight decrease in RMSE and a increase in R^2 , which is also expected because the model is trained with more data. Overall, we can conclude that 2d gives a better performance than 2c.

2c ->

```
In [208]: model7 = LinearRegression()
data4 = {}
model7.fit(X_train, Y_train)
data4["train-rmse"] = sqrt(mean_squared_error(Y_train, model7.predict(X_train)))
data4["train-r2"] = metrics.r2_score(Y_train, model7.predict(X_train))
data4["val-rmse"] = sqrt(mean_squared_error(Y_val, model7.predict(X_val)))
data4["val-r2"] = metrics.r2_score(Y_val, model7.predict(X_val))
data4["test-rmse"] = sqrt(mean_squared_error(Y_test, model7.predict(X_test)))
data4["test-r2"] = metrics.r2_score(Y_test, model7.predict(X_test))
data4

Out[208]: {'train-rmse': 94.16443818232597,
'train-r2': 0.2526942965780995,
'val-rmse': 91.61037186132472,
'val-r2': 0.12024864421105896,
'test-rmse': 88.04916726513582,
'test-r2': 0.06119046785291704}
```

2d->

```
In [212]: comb_datax = np.concatenate((X_train, X_val))
comb_datay = np.concatenate((Y_train, Y_val))
model = LinearRegression()
data3 = {}
model.fit(comb_datax, comb_datay)
data3["train-rmse"] = sqrt(mean_squared_error(Y_train, model.predict(X_train)))
data3["train-r2"] = metrics.r2_score(Y_train, model.predict(X_train))
data3["val-rmse"] = sqrt(mean_squared_error(Y_val, model.predict(X_val)))
data3["val-r2"] = metrics.r2_score(Y_val, model.predict(X_val))
data3["test-rmse"] = sqrt(mean_squared_error(Y_test, model.predict(X_test)))
data3["test-r2"] = metrics.r2_score(Y_test, model.predict(X_test))
data3

Out[212]: {'train-rmse': 95.02148918904848,
'train-r2': 0.2390289718836781,
'val-rmse': 85.00017194193263,
'val-r2': 0.2426262282210745,
'test-rmse': 82.69632325248155,
'test-r2': 0.17186834357698244}
```

2h) Comparing RMSE and R^2 for different values of λ

λ	Train RMSE	Train R^2	Val RMSE	Val R^2	Test RMSE	Test R^2
1010	95.358	0.23362	88.71337	0.1750096656	82.95	0.16677
1015	95.363	0.23354	88.71336	0.175009877	82.947	0.16682
1020	95.368	0.2334	88.71337	0.1750096655	82.9450	0.16687

These are the values I tested λ on.

From the table we can see that there is a decrease in Train R^2 whereas Val R^2 and Test R^2 increased for increase in λ . I chose the λ because of sharp increase in Val R^2 at 1015 and a slight decrease in later value. This data is for Ridge regression.

```
In [217]: model8 = Ridge(alpha=1010)
model8.fit(X_train, Y_train)
data5 = {}
data5["train-rmse"] = sqrt(mean_squared_error(Y_train, model8.predict(X_train)))
data5["train-r2"] = metrics.r2_score(Y_train, model8.predict(X_train))
data5["val-rmse"] = sqrt(mean_squared_error(Y_val, model8.predict(X_val)))
data5["val-r2"] = metrics.r2_score(Y_val, model8.predict(X_val))
data5["test-rmse"] = sqrt(mean_squared_error(Y_test, model8.predict(X_test)))
data5["test-r2"] = metrics.r2_score(Y_test, model8.predict(X_test))
data5

Out[217]: {'train-rmse': 95.36813338271575,
'train-r2': 0.23346670707893213,
'val-rmse': 88.7133754131302,
'val-r2': 0.17500966552596675,
'test-rmse': 82.94507166042648,
'test-r2': 0.1668788533804081}
```

```
In [219]: model8 = Ridge(alpha=1015)
model8.fit(X_train, Y_train)
data5 = {}
data5["train-rmse"] = sqrt(mean_squared_error(Y_train, model8.predict(X_train)))
data5["train-r2"] = metrics.r2_score(Y_train, model8.predict(X_train))
data5["val-rmse"] = sqrt(mean_squared_error(Y_val, model8.predict(X_val)))
data5["val-r2"] = metrics.r2_score(Y_val, model8.predict(X_val))
data5["test-rmse"] = sqrt(mean_squared_error(Y_test, model8.predict(X_test)))
data5["test-r2"] = metrics.r2_score(Y_test, model8.predict(X_test))
data5

Out[219]: {'train-rmse': 95.36335574564467,
'train-r2': 0.23354350686548375,
'val-rmse': 88.71336400061998,
'val-r2': 0.1750098777873078,
'test-rmse': 82.9475971322968,
'test-r2': 0.16682811965966005}
```

```
In [220]: model8 = Ridge(alpha=1020)
model8.fit(X_train, Y_train)
data5 = {}
data5["train-rmse"] = sqrt(mean_squared_error(Y_train, model8.predict(X_train)))
data5["train-r2"] = metrics.r2_score(Y_train, model8.predict(X_train))
data5["val-rmse"] = sqrt(mean_squared_error(Y_val, model8.predict(X_val)))
data5["val-r2"] = metrics.r2_score(Y_val, model8.predict(X_val))
data5["test-rmse"] = sqrt(mean_squared_error(Y_test, model8.predict(X_test)))
data5["test-r2"] = metrics.r2_score(Y_test, model8.predict(X_test))
data5

Out[220]: {'train-rmse': 95.36813338271575,
'train-r2': 0.23346670707893213,
'val-rmse': 88.7133754131302,
'val-r2': 0.17500966552596675,
'test-rmse': 82.94507166042648,
'test-r2': 0.1668788533804081}
```

For Lasso, I chose 7 as the R^2 values kept decreasing as the value of λ increased later on.

λ	Train RMSE	Train R^2	Val RMSE	Val R^2	Test RMSE	Test R^2
1	94.95	0.240117	89.7719	0.15520	85.25414	0.119847
10	98.677	0.1793	89.670	0.1571134	83.67966	0.15205
7	98.050	0.189735	89.0653	0.168450	84.0079	0.14539

```
In [227]: model9 = Lasso(alpha = 1)
model9.fit(X_train, Y_train)
data6 = {}
data6["train-rmse"] = sqrt(mean_squared_error(Y_train, model9.predict(X_train)))
data6["train-r2"] = metrics.r2_score(Y_train, model9.predict(X_train))
data6["val-rmse"] = sqrt(mean_squared_error(Y_val, model9.predict(X_val)))
data6["val-r2"] = metrics.r2_score(Y_val, model9.predict(X_val))
data6["test-rmse"] = sqrt(mean_squared_error(Y_test, model9.predict(X_test)))
data6["test-r2"] = metrics.r2_score(Y_test, model9.predict(X_test))
data6

Out[227]: {'train-rmse': 98.0507919899548,
'train-r2': 0.18973576513962997,
'val-rmse': 89.06535228802454,
'val-r2': 0.16845025854750784,
'test-rmse': 84.00794179055323,
'test-r2': 0.14539058353012047}
```

```
In [228]: model9 = Lasso(alpha = 10)
model9.fit(X_train, Y_train)
data6 = {}
data6["train-rmse"] = sqrt(mean_squared_error(Y_train, model9.predict(X_train)))
data6["train-r2"] = metrics.r2_score(Y_train, model9.predict(X_train))
data6["val-rmse"] = sqrt(mean_squared_error(Y_val, model9.predict(X_val)))
data6["val-r2"] = metrics.r2_score(Y_val, model9.predict(X_val))
data6["test-rmse"] = sqrt(mean_squared_error(Y_test, model9.predict(X_test)))
data6["test-r2"] = metrics.r2_score(Y_test, model9.predict(X_test))
data6

Out[228]: {'train-rmse': 98.67700162145417,
'train-r2': 0.17935307431419267,
'val-rmse': 89.6704241982925,
'val-r2': 0.15711349376385408,
'test-rmse': 83.67966564890754,
'test-r2': 0.1520566136751752}
```

```
In [229]: model9 = Lasso(alpha = 7)
model9.fit(X_train, Y_train)
data6 = {}
data6["train-rmse"] = sqrt(mean_squared_error(Y_train, model9.predict(X_train)))
data6["train-r2"] = metrics.r2_score(Y_train, model9.predict(X_train))
data6["val-rmse"] = sqrt(mean_squared_error(Y_val, model9.predict(X_val)))
data6["val-r2"] = metrics.r2_score(Y_val, model9.predict(X_val))
data6["test-rmse"] = sqrt(mean_squared_error(Y_test, model9.predict(X_test)))
data6["test-r2"] = metrics.r2_score(Y_test, model9.predict(X_test))
data6

Out[229]: {'train-rmse': 98.0507919899548,
'train-r2': 0.18973576513962997,
'val-rmse': 89.06535228802454,
'val-r2': 0.16845025854750784,
'test-rmse': 84.00794179055323,
'test-r2': 0.14539058353012047}
```

2j) Comparing RMSE and R^2 for combinedly trained data with optimal parameters

Model	λ	Train RMSE	Train R^2	Val RMSE	Val R^2	Test RMSE	Test R^2
Ridge	1015	95.599	0.22973	85.405	0.23538	81.5823	0.194029
Lasso	7	98.3208	0.18526	88.7759	0.17384	84.2093	0.14128

Here, we used the optimal parameter λ from the previous question 2h and applied it in 2i. The values are produced above.

```
In [231]: comb_datax = np.concatenate((X_train, X_val))
comb_datay = np.concatenate((Y_train, Y_val))
model = Ridge(alpha = 1015)
model.fit(comb_datax, comb_datay)
data1 = {}
data1["train-rmse"] = sqrt(mean_squared_error(Y_train, model.predict(X_train)))
data1["train-r2"] = metrics.r2_score(Y_train, model.predict(X_train))
data1["val-rmse"] = sqrt(mean_squared_error(Y_val, model.predict(X_val)))
data1["val-r2"] = metrics.r2_score(Y_val, model.predict(X_val))
data1["test-rmse"] = sqrt(mean_squared_error(Y_test, model.predict(X_test)))
data1["test-r2"] = metrics.r2_score(Y_test, model.predict(X_test))
data1
```

```
Out[231]: {'train-rmse': 95.59971064405147,
'train-r2': 0.22973952493994287,
'val-rmse': 85.40540965779748,
'val-r2': 0.2353874657286551,
'test-rmse': 81.58233146326182,
'test-r2': 0.19402938228180278}
```

```
In [230]: comb_datax = np.concatenate((X_train, X_val))
comb_datay = np.concatenate((Y_train, Y_val))
model = Lasso(alpha = 7)
model.fit(comb_datax, comb_datay)
data1 = {}
data1["train-rmse"] = sqrt(mean_squared_error(Y_train, model.predict(X_train)))
data1["train-r2"] = metrics.r2_score(Y_train, model.predict(X_train))
data1["val-rmse"] = sqrt(mean_squared_error(Y_val, model.predict(X_val)))
data1["val-r2"] = metrics.r2_score(Y_val, model.predict(X_val))
data1["test-rmse"] = sqrt(mean_squared_error(Y_test, model.predict(X_test)))
data1["test-r2"] = metrics.r2_score(Y_test, model.predict(X_test))
data1
```


```
Out[230]: {'train-rmse': 98.32087385322708,
'train-r2': 0.1852658559201099,
'val-rmse': 88.7759830225751,
'val-r2': 0.173844815285358,
'test-rmse': 84.20939921311948,
'test-r2': 0.14128683233875394}
```

For Ridge : There is a slight increase in RMSE of Train whereas it decreased in case of both Validation and Test (from 2h). Regarding R^2 for train it has decreased slightly whereas for Validation it went from 0.175009 to 0.23538 and for the test it increased from 0.16682 to 0.194029.

For Lasso : There is a slight increase in RMSE of Train whereas it decreased in case of both Validation and Test for the Lasso too (from 2h). Regarding R^2 for train it has decreased slightly whereas for Validation it went from 0.168450 to 0.17384 and for the test it decreased from 0.14539 to 0.14128.

So we can infer from the above information that when the model is trained with more data, it presents us with more accuracy in all the above three cases.

2k) Given coefficient path plots for both ridge and lasso :




We can see a dotted line which represents the optimal value for both lasso and ridge.

2l) Three key observations that we can pick from the above plots :

1. The coefficients vary smoothly along the regularization path. As lambda increases, the coefficient values move towards zero. This shows the regularization is working as expected.
2. The lasso model sets more coefficients exactly to zero compared to ridge.
3. The optimal lambda is smaller for lasso than ridge, corresponding to less overall regularization strength.

3f) From the optimization parameter of ridge(1015), some of the good learning rates are 0.000001 and 0.000001.





From the chosen learning rates, by observation I chose the learning rate of 0.0000001 to be optimal as Objective decreases smoothly over epochs and it does not diverge or become unstable.

The RMSE values for Train, Validation and Test data are 105.81903171862118, 96.59776672570594, 90.0388759165188.

The R^2 values for Train, Validation and Test data are 0.05626073596552117, 0.021851429325226768, 0.01828118710364468

```
In [69]: elas_ridge1 = ElasticNet(1015, 0.5, 0.0000001, 100, 100)
dict2 = elas_ridge1.train(X_train, Y_train)
plt.plot(dict2.keys(), dict2.values(), label = 'LR = 0.0000001')
plt.xlabel('Epoch')
plt.ylabel('Loss ')
plt.title("For learning rate of 0.0000001")
```

From the optimization parameter of lasso (7), some of the good learning rates are 0.00001 and 0.000001.



From the chosen learning rates, by observation I chose the learning rate of 0.000001 to be optimal as Objective decreases smoothly over epochs and it does not diverge or become unstable. The RMSE values for train, validation and test data are 98.41135822347093, 89.74285386580931, 83.65010817548402 and The R^2 Values for Train, Validation and Test data are 0.18376557168723273, 0.15575129109868457, 0.15265553188314407.

```
In [55]: import matplotlib.pyplot as plt
elas_lasso = ElasticNet(9, 0.5, 0.000001, 100, 1000)
dic = elas_lasso.train(X_train, Y_train)
plt.plot(dic.keys(), dic.values())
plt.xlabel('Epoch')
plt.ylabel('Loss ')
plt.title("For learning rate of 0.000001")
```


3g) For Ridge λ :

Alpha s	Train R^2	Validation R^2	Test R^2	Train RMSE	Validation RMSE	Test RMSE
0.0	0.0562607521910835	0.021851438829579184	0.018281196451523707	105.8190308089561	96.59776625640136	90.03887548784589
0.11	0.05626074858540353	0.02185143671750156	0.018281194374218046	105.81903101110387	96.59776636069122	90.0388755831065
0.22	0.05626074497972311	0.021851434605423492	0.01828119229691183	105.81903121325166	96.59776646498112	90.03887567836713
0.33	0.056260741374042356	0.02185143249334509	0.01828119021960528	105.81903141539946	96.59776656927103	90.03887577362778
0.44	0.05626073776836171	0.021851430381266357	0.018281188142298066	105.81903161754727	96.59776667356097	90.0388758688847
0.55	0.05626073416268074	0.0218514282691874	0.01828118606499063	105.81903181969508	96.59776677785091	90.03887596414916
0.66	0.05626073055699943	0.02185142615710789	0.018281183987682748	105.81903202184292	96.59776688214087	90.03887605940987
0.77	0.056260726951317785	0.021851424045027934	0.01828118191037431	105.81903222399079	96.59776698643086	90.0388761546706
0.88	0.05626072334563592	0.021851421932947868	0.018281179833065653	105.81903242613865	96.59776709072085	90.03887624993136
1.0	0.05626071973995406	0.021851419820867468	0.018281177755756328	105.81903262828652	96.59776719501086	90.03887634519214

```

]: alpha = np.linspace(0,1,10)
for i in alpha :
    elas_ridge = ElasticNet(1015, i ,0.0000001, 100, 100)
    elas_ridge.train(X_train, Y_train)
    b = elas_ridge.predict(X_train)
    b_val = elas_ridge.predict(X_val)
    b_test = elas_ridge.predict(X_test)
    r2_train = r2_score(Y_train, b)
    r2_val = r2_score(Y_val, b_val)
    r2_test = r2_score(Y_test, b_test)
    print ("The r2 values for Train, Validation and Test datas are ", r2_train, r2_val, r2_test)
    rmse = sqrt(mean_squared_error(Y_train,b))
    rmse_val = sqrt(mean_squared_error(Y_val,b_val))
    rmse_test = sqrt(mean_squared_error(Y_test,b_test))
    print ("The rmse values for Train, Validation and Test datas are ", rmse, rmse_val, rmse_test)

```

For Lasso λ :

Alpha s	Train R^2	Validation R^2	Test R^2	Train RMSE	Validation RMSE	Test RMSE
0.0	0.12719434498203 13	0.0789545360304804 2	0.071990374087808 4	101.76455712717 384	93.735738996878 15	87.5412509221681 3
0.11	0.12719434451815 503	0.0789545355310341 5	0.071990373595943 62	101.76455715421 662	93.735739022292 75	87.5412509453675
0.22	0.12719434405427 865	0.0789545350315877 8	0.071990373104078 85	101.76455718125 939	93.735739047707 33	87.5412509685668 6
0.33	0.12719434359040 25	0.0789545345321414 1	0.071990372612213 97	101.76455720830 214	93.735739073121 9	87.5412509917662 2
0.44	0.12719434312652 6	0.0789545340326951 5	0.071990372120348 98	101.76455723534 492	93.735739098536 5	87.5412510149655 8
0.55	0.12719434266264 973	0.0789545335332488 8	0.071990371628484 2	101.76455726238 77	93.735739123951 06	87.5412510381649 4
0.66	0.12719434219877 335	0.0789545330338025 1	0.071990371136619 44	101.76455728943 047	93.735739149365 66	87.5412510613643
0.77	0.12719434173489 708	0.0789545325343563 6	0.071990370644754 66	101.76455731647 324	93.735739174780 22	87.5412510845636 7
0.88	0.12719434127102 092	0.0789545320349099 9	0.071990370152889 9	101.76455734351 6	93.735739200194 8	87.5412511077630 2
1.0	0.12719434080714 453	0.0789545315354638 3	0.071990369661025 12	101.76455737055 879	93.735739225609 38	87.5412511309623 8

```
In [71]: alphas = np.linspace(0,1,10)
for i in alphas :
    elas_lasso2 = ElasticNet(9, i ,0.000001, 100, 100)
    elas_lasso2.train(X_train, Y_train)
    c = elas_lasso2.predict(X_train)
    c_val = elas_lasso2.predict(X_val)
    c_test = elas_lasso2.predict(X_test)
    r2_train1 = r2_score(Y_train, c)
    r2_val1 = r2_score(Y_val, c_val)
    r2_test1 = r2_score(Y_test, c_test)
    print ("The r2 values for Train, Validation and Test datas are ", r2_train1, r2_val1, r2_test1)
    rmse1 = sqrt(mean_squared_error(Y_train,c))
    rmse_val1 = sqrt(mean_squared_error(Y_val,c_val))
    rmse_test1= sqrt(mean_squared_error(Y_test,c_test))
    print ("The rmse values for Train, Validation and Test datas are ", rmse1, rmse_val1, rmse_test1)
```

I have tested for both values of λ For Ridge and lasso, we can see when alpha value is near zero, it shows some possible lasso model behavior and when it increases towards 1, it tends towards ridge model.

3h) From my observations in 2h, 2j, when compared different values of λ to find the optimal parameter based on validation data performance. In 3g, we test models in a range of alphas(0,1) . From my observation, when we used optimal parameter and different λ values, we got better RMSE and R^2 compared to SGD variants. We can infer from the results that when we train the model using training and validation data with an optimal parameter, it showed good results. From both lasso and ridge models, Ridge displayed a better performance than lasso when trained with both the data of train and validation.

Differences between SGD variants of Ridge and Lasso :

- SGD variants use stochastic gradient descent to optimize the regression coefficients, while standard implementations use closed-form solutions (for Ridge) or coordinate descent (for LASSO).
- For LASSO, the SGD variant typically uses a constant learning rate. So standard LASSO can converge faster and more precisely.
- SGD estimators have more hyperparameters to tune (learning rate, batch size, iterations) compared to the closed-form solutions.

3j)

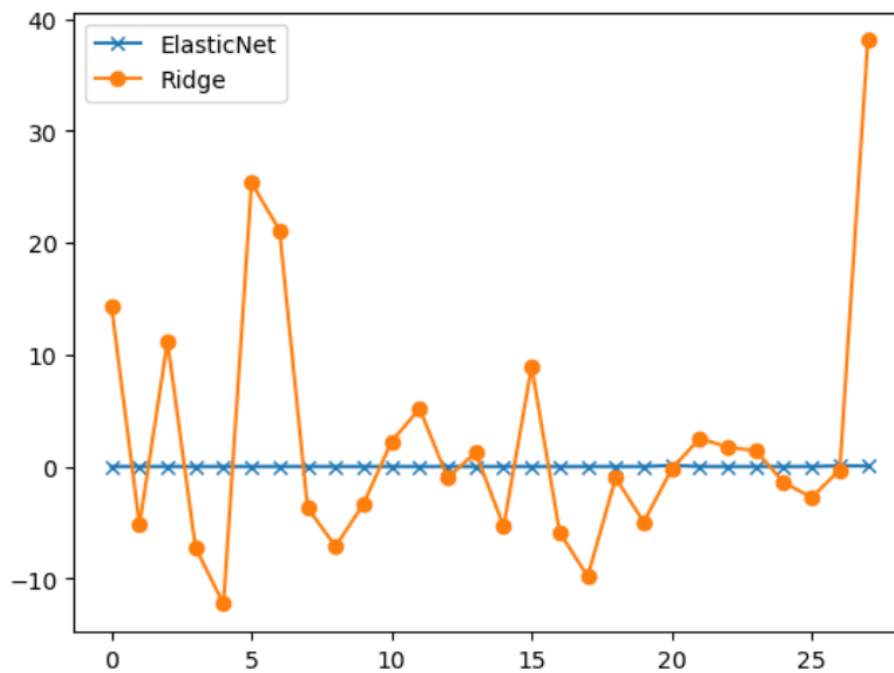
For final coefficients for the best Test R^2 in elastic net model are given below :

```
array([[0.00475841],
       [0.00322165],
       [0.00627088],
       [0.00322873],
       [0.00530439],
       [0.00336598],
       [0.00605122],
       [0.00309992],
       [0.00553556],
       [0.00271043],
       [0.00712199],
       [0.0012842 ],
       [0.00836365],
       [0.00290355],
       [0.00428036],
       [0.00326167],
       [0.00492464],
       [0.00270387],
       [0.0054718 ],
       [0.00113735],
       [0.1109303 ],
       [0.0091522 ],
       [0.00088932],
       [0.00540959],
       [0.00014211],
       [0.00476814],
       [0.07411704],
       [0.0504452 ]])
```

The coefficients of Ridge without SGD that performed well on validation data are :

```
array([ 14.36479005, -5.22738175, 11.18651141, -7.250327 ,  
       -12.27338396, 25.36479653, 21.11861461, -3.69370588,  
        -7.13401571, -3.33222394,  2.25977857,  5.16176822,  
        -0.98252213,  1.23666284, -5.35249052,  8.95058434,  
        -5.89212568, -9.82002434, -1.00787877, -4.96571115,  
        -0.234283  ,  2.53411894,  1.72552405,  1.42319065,  
        -1.41260902, -2.76949532, -0.29383492, 38.08538165])
```

The comparison between the coefficients can be plotted as :



From the above we can say the coefficients different in most of the time but they are equal some times.